



S-109A Introduction to Data Science

Homework 1

Harvard University

Summer 2019

Instructors: Pavlos Protopapas and Kevin Rader

Main Theme: Data Collection - Web Scraping - Data Parsing

Learning Objectives

In this homework, your goal is to learn how to acquire, parse, clean, and analyze data. Initially you read the data from a file, then you scrape them directly from a website. You look for specific pieces of information by parsing the data, you clean the data to prepare them for analysis, and finally, you answer some questions.

Instructions

- To submit your assignment follow the instructions given in Classroom.
- The deliverables in Classroom are: a) This python notebook with your code and answers, b) a .pdf version of this notebook, c) The BibTex file you created. d) The JSON file you created.
- Exercise **responsible scraping**. Web servers can become slow or unresponsive if they receive too many requests from the same source in a short amount of time. Use a delay of 10 seconds between requests in your code. This helps not to get blocked by the target website. Run the webpage fetching part of the homework only once and do not re-run after you have saved the results in the JSON file (details below).
- Web scraping requests can take several minutes. This is another reason why you should not wait until the last minute to do this homework.

In [1]:

```
1  # import the necessary libraries
2  %matplotlib inline
3  import numpy as np
4  import scipy as sp
5  import matplotlib as mpl
6  import matplotlib.cm as cm
7  import matplotlib.pyplot as plt
8  import pandas as pd
9  import time
10 pd.set_option('display.width', 500)
11 pd.set_option('display.max_columns', 100)
12 pd.set_option('display.notebook_repr_html', True)
13
```

Part A [50 pts]: Help a professor convert his publications to bibTex

Overview

In Part 1 your goal is to parse the HTML page of a Professor containing some of his publications, and answer some questions. This page is provided to you in the file `data/publist_super_clean.html`. There are 44 publications in descending order from No. 244 to No. 200.

You are to use python's **regular expressions**, a powerful way of parsing text. You may **not** use any parsing tool such as Beautiful Soup yet. In doing so you will get more familiar with three of the common file formats for storing and transferring data, which are:

- CSV, a text-based file format used for storing tabular data that are separated by some delimiter, usually comma or space.
- HTML/XML, the stuff the web is made of.
- JavaScript Object Notation(JSON), a text-based open standard designed for transmitting structured data over the web.

Question 1: Parsing using Regular Expressions

1.1 Write a function called `get_pubs` that takes a `.html` filename as an input and returns a string containing the HTML page in this file (see definition below). Call this function using `data/publist_super_clean.html` as input and name the returned string `prof_pubs`.

1.2 Calculate how many times the author named 'C.M. Friend' appears in the list of publications.

1.3 Find all unique journals and copy them in a variable named `journals`.

1.4 Create a list named `pub_authors` whose elements are strings containing the authors' names for each paper.

Hints

- Look for patterns in the HTML tags that reveal where each piece of information such as the title of the paper, the names of the authors, the journal name, is stored. For example, you might notice that the journal name(s) is contained between the `<I>` HTML tag.
- Each publication has multiple authors.
- C.M. Friend also shows up as Cynthia M. Friend in the file. Count just C. M. Friend.
- There is a comma at the end of the string of authors. You can choose to keep it in the string or remove it and put it back when you write the string as a BibTex entry.
- You want to remove duplicates from the list of journals.

Resources

- **Regular expressions:** a) <https://docs.python.org/3.3/library/re.html> (<https://docs.python.org/3.3/library/re.html>), b) <https://regexone.com> (<https://regexone.com>), and c) <https://docs.python.org/3/howto/regex.html> (<https://docs.python.org/3/howto/regex.html>).
- **** HTML:**** if you are not familiar with HTML see <https://www.w3schools.com/html/> (<https://www.w3schools.com/html/>) or one of the many tutorials on the internet.
- **** Document Object Model (DOM):**** for more on this programming interface for HTML and XML documents see https://www.w3schools.com/js/js_htmldom.asp (https://www.w3schools.com/js/js_htmldom.asp).

1.1

In [2]:

```
1 # import the regular expressions library
2 import re
```

In [3]:

```
1 # use this file
2 pub_filename = 'data/publist_super_clean.html'
```

In [4]:

```
1 # definition of get_pubs
2 def get_pubs(filename: str) -> str:
3
4     '''Open the file using the filename.
5
6     Args:
7         filename: A string name of the file.
8
9     Returns:
10         A string containing the HTML page ready to be parsed.
11     '''
12
13     f = open(filename, 'r')
14     profpubs = f.read()
15     return profpubs
```

In [5]:

```
1 # check your code
2 # print (prof_pubs)
3 prof_pubs = get_pubs(pub_filename)
4 print(prof_pubs)
```

```
<!DOCTYPE HTML PUBLIC "-//W3C//DTD HTML 4.01 Transitional//EN"
"http://www.w3.org/TR/html4/loose.dtd">
```

```
<TITLE>Kaxiras E journal publications</TITLE>
```

```
<HEAD>
```

```
<meta http-equiv="Content-Type" content="text/html; charset=UTF-8">
```

```
<LINK REL="stylesheet" TYPE="text/css" HREF="../styles/style_pubs.cs
s">
```

```
<META NAME="description" CONTENT="">
```

```
<META NAME="keywords" CONTENT="Kaxiras E, Multiscale Methods, Comput
ational Materials" >
```

```
</HEAD>
```

```
<BODY>
```

```
<OL START=244>
```

```
<LI>
```

```
<A HREF="Papers/2011/PhysRevB_84_125411_2011.pdf" target="paper244">
"Approaching the intrinsic band gap in suspended high-mobility
```

```
graphene nanoribbons"</A>
```

You should see an HTML page

```

<LI>
<A HREF="Papers/2011/PhysRevB_84_125411_2011.pdf" target="paper244">
"Approaching the intrinsic band gap in suspended high-mobility graphen
e nanoribbons"</A>
<BR>Ming-Wei Lin, Cheng Ling, Luis A. Agapito, Nicholas Kioussis, Yiyang Zh
ang, Mark Ming-Cheng Cheng,
<I>PHYSICAL REVIEW B </I> <b>84</b>, 125411 (2011)
<BR>
</LI>
</OL>

<OL START=243>
<LI>
<A HREF="Papers/2011/PhysRevB_84_035325_2011.pdf" target="paper243">
"Effect of symmetry breaking on the optical absorption of semiconducto
r nanoparticles"</A>
<BR>JAdam Gali, Efthimios Kaxiras, Gergely T. Zimanyi, Sheng Meng,
<I>PHYSICAL REVIEW B </I> <b>84</b>, 035325 (2011)
<BR>
</LI>
</OL>

<OL START=242>
<LI>
<A HREF="Papers/2011/PhysRevB_83_054204_2011.pdf" target="paper242">
"Influence of CH2 content and network defects on the elastic propertie
s of organosilicate glasses"</A>
<BR>Jan M. Knaup, Han Li, Joost J. Vlassak, and Efthimios Kaxiras,
<I>PHYSICAL REVIEW B </I> <b>83</b>, 054204 (2011)
<BR>
</LI>
</OL>

```

1.2

In [6]:

```

1 # your code here
2 print(prof_pubs.count('C.M. Friend'))

```

5

1.3

In [7]:

```

1 # your code here
2 journals = re.findall('<I>(.*?)</I>', prof_pubs)
3 journals = sorted(list(set(journals)))

```

In [8]:

```
1 # check your code: print journals
2 journals
```

Out[8]:

```
['2010 ACM/IEEE International Conference for High Performance ',
'ACSNano. ',
'Ab initio',
'Acta Mater. ',
'Catal. Sci. Technol. ',
'Chem. Eur. J. ',
'Comp. Phys. Comm. ',
'Concurrency Computat.: Pract. Exper. ',
'Energy & Environmental Sci. ',
'Int. J. Cardiovasc. Imaging ',
'J. Chem. Phys. ',
'J. Chem. Theory Comput. ',
'J. Phys. Chem. B ',
'J. Phys. Chem. C ',
'J. Phys. Chem. Lett. ',
'J. Stat. Mech: Th. and Exper. ',
'Langmuir ',
'Molec. Phys. ']
```

Your output should look like this (remember, no duplicates):

```
'ACSNano.',
'Ab initio',
'Ab-initio',
'Acta Mater.',
'Acta Materialia',
'Appl. Phys. Lett.',
'Applied Surface Science',
'Biophysical J.',
'Biosensing Using Nanomaterials',

...

'Solid State Physics',
'Superlattices and Microstructures',
'Surf. Sci.',
'Surf. Sci. Lett.',
'Surface Science',
'Surface Review and Letters',
'Surface Sci. Lett.',
'Surface Science Lett.',
'Thin Solid Films',
'Top. Catal.',
'Z']
```

In [9]:

```
1 # our code here
2 pub_authors = re.findall('<BR> ?(.+)', prof_pubs)
3 pub_authors = sorted(list(set(pub_authors)))
```

In [10]:

```
1 # check your code: print the list of strings containing the author(s)' names
2 for item in pub_authors:
3     print (item)
```

```
A. Gali and E. Kaxiras,
A. Gali, E. Janzen, P. Deak, G. Kresse and E. Kaxiras,
A. Peters, S. Melchionna, E. Kaxiras, J. Latt, J. Sircar, S. Succi,
Bingjun Xu, Jan Haubrich, Thomas A. Baker, Efthimios Kaxiras, and Cy
nthia M. Friend,
C.E. Lekka, J. Ren, S. Meng and E. Kaxiras,
C.L. Chang, S.K.R.S. Sankaranarayanan, D. Ruzmetov, M.H. Engelhard,
E. Kaxiras and S. Ramanathan,
E. Kaxiras and S. Succi,
E. Manousakis, J. Ren, S. Meng and E. Kaxiras,
E.M. Kotsalis, J.H. Walther, E. Kaxiras and P. Koumoutsakos,
F.J. Rybicki, S. Melchionna, D. Mitsouras, A.U. Coskun, A.G. Whitmor
e, E. Kaxiras, S. Succi, P.H. Stone and C.L. Feldman,
H. Chen, W.G. Zhu, E. Kaxiras, and Z.Y. Zhang,
H. Li, J.M. Knaup, E. Kaxiras and J.J. Vlassak,
H.P. Chen, R.K. Kalia, E. Kaxiras, G. Lu, A. Nakano, K. Nomura,
J R Maze, A Gali, E Togan, Y Chu, A Trifonov,
J. Ren, E. Kaxiras and S. Meng,
JAdam Gali, Efthimios Kaxiras, Gergely T. Zimanyi, Sheng Meng,
Jan Haubrich, Efthimios Kaxiras, and Cynthia M. Friend
```

Your output should look like this (a line for each paper's author(s) string, with or without the comma)

```
S. Meng and E. Kaxiras,
G. Lu and E. Kaxiras,
E. Kaxiras and S. Yip,
...
Simone Melchionna, Efthimios Kaxiras, Massimo Bernaschi and Sauro Succi,
J R Maze, A Gali, E Togan, Y Chu, A Trifonov,
E Kaxiras, and M D Lukin,
```

Question 2: Parsing and Converting to bibTex using Beautiful Soup

A lot of the bibliographic and publication information is displayed in various websites in a not-so-structured HTML files. Some publishers prefer to store and transmit this information in a .bibTex file which has the following format:

```

@article { _number_
  author = John Doyle
  title = Interaction between atoms
  URL = Papers/PhysRevB_81_085406_2010.pdf
  journal = Phys. Rev. B
  volume = 81
}

@article
{
  author = Ming-Wei Lin, Cheng Ling, Luis A. Agapito, Nicholas Kioussis,
  Yiyang Zhang, Mark Ming-Cheng Cheng
  title = "Approaching the intrinsic band gap in suspended high-mobility
  graphene nanoribbons"
  URL = Papers/2011/PhysRevB_84_125411_2011.pdf
  journal = PHYSICAL REVIEW B
  volume = 84
}

```

About the [bibTex format \(http://www.bibtex.org\)](http://www.bibtex.org).

In Question 2 you are given an .html file containing a list of papers scraped from the author's website and you are to write the information into .bibTex format. We used regular expressions for parsing HTML in the previous question but just regular expressions are hard to use in parsing real-life websites. A useful tool is [BeautifulSoup] (<http://www.crummy.com/software/BeautifulSoup/> (<http://www.crummy.com/software/BeautifulSoup/>)) (BS). You will parse the same file, this time using BS, which makes parsing HTML a lot easier.

2.1 Write a function called `make_soup` that accepts a filename for an HTML file and returns a BS object.

2.2 Write a function that reads in the BS object, parses it, converts it into the .bibTex format using python string manipulation and regular expressions, and writes the data into `publist.bib`. You will need to create that file in your folder.

HINT

- Inspect the HTML code for tags that indicate information chunks such as `title` of the paper. You had already done this in Part 1 when you figured out how to get the name of the journal from the HTML code. The `find_all` method of BeautifulSoup might be useful.
- Question 2.2 is better handled if you break the code into functions, each performing a small task such as finding the author(s) for each paper.
- Make sure you catch exceptions when needed.
- Regular expressions are a great tool for string manipulation.

Resources

- [BeautifulSoup Tutorial \(https://www.dataquest.io/blog/web-scraping-tutorial-python/\)](https://www.dataquest.io/blog/web-scraping-tutorial-python/).
- More about the [BibTex format \(http://www.bibtex.org\)](http://www.bibtex.org).

In [11]:

```
1 # import the necessary libraries
2 from bs4 import BeautifulSoup
3 from sys import argv
4 from urllib.request import urlopen
5 from urllib.error import HTTPError
```

2.1

In [12]:

```
1 # your code here
2
3 # definition of make_soup
4 def make_soup(filename: str) -> BeautifulSoup:
5     '''Open the file and convert into a BS object.
6
7     Args:
8         filename: A string name of the file.
9
10    Returns:
11        A BS object containing the HTML page.
12    ...
13    return BeautifulSoup(get_pubs(filename))
14
15 soup = make_soup(pub_filename)
```

In [13]:

```
1 # check your code: print the BeautifulSoup object, you should see an HTML page
2 print(soup)
```

```
<!DOCTYPE HTML PUBLIC "-//W3C//DTD HTML 4.01 Transitional//EN" "http://www.w3.org/TR/html4/loose.dtd">
<html><head><title>Kaxiras E journal publications</title>
<meta content="text/html; charset=utf-8" http-equiv="Content-Type"/>
<link href="../styles/style_pubs.css" rel="stylesheet" type="text/css"/>
<meta content="" name="description"/>
<meta content="Kaxiras E, Multiscale Methods, Computational Materials" name="keywords"/>
</head><body>
<ol start="244">
<li>
<a href="Papers/2011/PhysRevB_84_125411_2011.pdf" target="paper244">
"Approaching the intrinsic band gap in suspended high-mobility graphene nanoribbons"</a>
<br/>Ming-Wei Lin, Cheng Ling, Luis A. Agapito, Nicholas Kioussis, Yiyang Zhang, Mark Ming-Cheng Cheng,
<i>PHYSICAL REVIEW B</i> <b>84</b>, 125411 (2011)
<br/>
```

Your output should look like this:


```
<!DOCTYPE HTML PUBLIC "-//W3C//DTD HTML 4.01 Transitional//EN"
"http://www.w3.org/TR/html4/loose.dtd">
```

```
<title>Kaxiras E journal publications</title>
<head>
<meta content="text/html; charset=utf-8" http-equiv="Content-Type"/>
<link href="../../styles/style_pubs.css" rel="stylesheet" type="text/css"/>
<meta content="" name="description"/>
<meta content="Kaxiras E, Multiscale Methods, Computational Materials" name
="keywords"/>
</head>
<body>
<ol start="244">
<li>
<a href="Papers/2011/PhysRevB_84_125411_2011.pdf" target="paper244">
"Approaching the intrinsic band gap in suspended high-mobility graphene nan
oribbons"</a>
<br/>Ming-Wei Lin, Cheng Ling, Luis A. Agapito, Nicholas Kioussis, Yiyang Z
hang, Mark Ming-Cheng Cheng,
<i>PHYSICAL REVIEW B </i> <b>84</b>, 125411 (2011)
<br/>
</li>
</ol>
<ol start="243">
<li>
<a href="Papers/2011/PhysRevB_84_035325_2011.pdf" target="paper243">
"Effect of symmetry breaking on the optical absorption of semiconductor nan
oparticles"</a>
<br/>JAdam Gali, Efthimios Kaxiras, Gergely T. Zimanyi, Sheng Meng,
<i>PHYSICAL REVIEW B </i> <b>84</b>, 035325 (2011)
<br/>
</li>
</ol>
...
```

In [14]:

```
1 # your code here
2 papers = soup.find_all('li')
3 bibtex = ''
4 for paper in papers:
5     bibtex += '@article\n'
6     bibtex += '{\n'
7     bibtex += 'author = '\n'
8     bibtex += re.findall(
9         '<br/>(.)+',
10         str(paper)
11     )[0].strip(' ').strip(',').strip(' ')
12     bibtex += '\n        title = '\n'
13     #print(paper)
14     bibtex += paper.find('a').getText().strip('\n').strip(' ')
15     bibtex += '\n        URL = '\n'
16     bibtex += paper.find('a')['href'].strip(' ')
17     bibtex += '\n        journal = '\n'
18     bibtex += paper.find('i').getText().strip(' ')
19     if paper.find('b') is not None:
20         bibtex += '\n        volume = '\n'
21         bibtex += paper.find('b').getText().strip(' ')
22     bibtex += '\n}\n\n'
23
24 with open('publist.bib', 'w') as f:
25     f.write(bibtex)
```

In [15]:

```
1 # check your code: print the BibTex file
2 f = open('publist.bib', 'r')
3 print (f.read())
```

```
@article
{
  author = Ming-Wei Lin, Cheng Ling, Luis A. Agapito, Nicholas Ki
oussis, Yiyang Zhang, Mark Ming-Cheng Cheng
  title = "Approaching the intrinsic band gap in suspended high-m
obility graphene nanoribbons"
  URL = Papers/2011/PhysRevB_84_125411_2011.pdf
  journal = PHYSICAL REVIEW B
  volume = 84
}

@article
{
  author = JAdam Gali, Efthimios Kaxiras, Gergely T. Zimanyi, She
ng Meng
  title = "Effect of symmetry breaking on the optical absorption
of semiconductor nanoparticles"
  URL = Papers/2011/PhysRevB_84_035325_2011.pdf
  journal = PHYSICAL REVIEW B
  volume = 84
}
```

Your output should look like this

```

@article
{
    author = Ming-Wei Lin, Cheng Ling, Luis A. Agapito, Nicholas Kioussis,
    Yiyang Zhang, Mark Ming-Cheng Cheng
    title = "Approaching the intrinsic band gap in suspended high-mobility
    graphene nanoribbons"
    URL = Papers/2011/PhysRevB_84_125411_2011.pdf
    journal = PHYSICAL REVIEW B
    volume = 84
}

...

@article
{
    author = E. Kaxiras and S. Succi
    title = "Multiscale simulations of complex systems: computation meets
    reality"
    URL = Papers/SciModSim_15_59_2008.pdf
    journal = Sci. Model. Simul.
    volume = 15
}

@article
{
    author = E. Manousakis, J. Ren, S. Meng and E. Kaxiras
    title = "Effective Hamiltonian for FeAs-based superconductors"
    URL = Papers/PhysRevB_78_205112_2008.pdf
    journal = Phys. Rev. B
    volume = 78
}

```

Part B [50 pts]: Follow the stars in IMDb's list of "The Top 100 Stars for 2017"

Overview

In Part 3 your goal is to extract information from IMDb's Top 100 Stars for 2017 (<https://www.imdb.com/list/ls025814950/>) and perform some analysis on each star in the list. In particular we are interested to know: a) how many performers made their first movie at 17? b) how many performers started as child actors? c) who is the most proliferate actress or actor in IMDb's list of the Top 100 Stars for 2017? . These questions are addressed in more details in the Questions below.

When data is **not** given to us in a file, we need to fetch them using one of the following ways:

- download a file from a source URL
- query a database
- query a web API
- scrape data from the web page

Question 1: Web Scraping Using Beautiful Soup

1.1 Download the webpage of the "Top 100 Stars for 2017" (<https://www.imdb.com/list/ls025814950/>) into a `requests` object and name it `my_page`. Explain what the following attributes are:

- `my_page.text`,
- `my_page.status_code`,
- `my_page.content`.

1.2 Create a BeautifulSoup object named `star_soup` giving `my_page` as input.

1.3 Write a function called `parse_stars` that accepts `star_soup` as its input and generates a list of dictionaries named `starlist` (see definition below). One of the fields of this dictionary is the `url` of each star's individual page, which you need to scrape and save the contents in the `page` field. Note that there is a ton of information about each star on these webpages.

1.4 Write a function called `create_star_table` to extract information about each star (see function definition for the exact information to extract). **Only extract information from the first box on each star's page. If the first box is acting, consider only acting credits and the star's acting debut, if the first box is Directing, consider only directing credits and directorial debut.**

1.5 Now that you have scraped all the info you need, it's a good practice to save the last data structure you created to disk. That way if you need to re-run from here, you don't need to redo all these requests and parsing. Save this information to a JSON file and **submit** this JSON file in Canvas with your notebook.

1.6 Import the contents of the teaching staff's JSON file (`data/staff_starinfo.json`) into a pandas dataframe. Check the types of variables in each column and clean these variables if needed. Add a new column to your dataframe with the age of each actor when they made first movie (name this column `age_at_first_movie`).

1.7 You are now ready to answer the following intriguing questions:

- How many performers made their first movie at 17?
- How many performers started as child actors? Define child actor as a person less than 12 years old.
- Who is the most prolific actress or actor in IMDb's list of the Top 100 Stars for 2017?

1.8 Make a plot of the number of credits versus the name of actor/actress.

Hints

- Create a variable that groups actors/actresses by the age of their first movie. Use pandas' `.groupby` to divide the dataframe into groups of performers that for example started performing as children (`age < 12`). The grouped variable is a `GroupBy` pandas object and this object has all of the information needed to then apply some operation to each of the groups.
- When cleaning the data make sure the variables with which you are performing calculations are in numerical format.
- The column with the year has some values that are double, e.g. `'2000-2001'` and the column with age has some empty cells. You need to deal with these before performing calculations on the data!
- You should include both movies and TV shows.

Resources

- The `requests` library makes working with HTTP requests powerful and easy. For more on the `requests` library see <http://docs.python-requests.org/> (<http://docs.python-requests.org/>).

In [16]:

```
1 import requests
```

1.1

In [17]:

```
1 # your code here
2 my_page = requests.get('https://www.imdb.com/list/ls025814950/')
```

In [18]:

```
1 print(my_page.text)
```

```
<!DOCTYPE html>
<html
  xmlns:og="http://ogp.me/ns#"
  xmlns:fb="http://www.facebook.com/2008/fbml">
  <head>

    <meta charset="utf-8">
    <meta http-equiv="X-UA-Compatible" content="IE=edge">

    <meta name="apple-itunes-app" content="app-id=342792525, app-argument=imdb:///list/ls025814950?src=mdot">

    <script type="text/javascript">var IMDbTimer={starttime: new
Date(), getTime()>
ntaliquellie</script>
```

In [30]:

```
1 print(my_page.status_code)
```

200

In [31]:

```
1 print(my_page.content)
```

```

b'\n\n\n\n<!DOCTYPE html>\n\n<html>\n\n    xmlns:og="http://ogp.me/ns#" \n
xmlns:fb="http://www.facebook.com/2008/fbml">\n        <head>\n
\n            <meta charset="utf-8">\n                <meta http-equiv="X-UA-Com
patible" content="IE=edge">\n\n                <meta name="apple-itunes-app" con
tent="app-id=342792525, app-argument=imdb:///list/ls025814950?src=md
ot">\n\n\n\n\n                <script type="text/javascript">var IMDbTimer={st
arttime: new Date().getTime(),pt:'\java\'};</script>\n\n\n<script>\n
if (typeof uet == '\function\') {\n                uet("bb", "LoadTitle", {wb:
1});\n            }\n</script>\n\n    <script>(function(t){ (t.events = t.events
|| {})[ "csm_head_pre_title" ] = new Date().getTime(); })(IMDbTimer);
</script>\n        <title>Top 100 Stars of 2017 - IMDb</title>\n\n    <s
cript>(function(t){ (t.events = t.events || {})[ "csm_head_post_titl
e" ] = new Date().getTime(); })(IMDbTimer);</script>\n\n<script>\n        i
f (typeof uet == '\function\') {\n                uet("be", "LoadTitle", {wb:
1});\n            }\n</script>\n\n<script>\n        if (typeof uex == '\function
\') {\n                uex("ld", "LoadTitle", {wb: 1});\n            }\n</script>\n\n\n
<link rel="canonical" href="https://www.imdb.com/list/ls025814950/"
/>\n            <meta property="og:url" content="http://www.imdb.com/lis
t/ls025814950/" />\n\n\n<script>\n        if (typeof uet == '\function\')

```

Your answers here:

- `my_page.content` : byte string of the HTML of the page;
- `my_page.text` : string of the HTML of the page;
- `my_page.status_code` : '200' means it was a successful request.

1.2

In [32]:

```
1 # your code here
2 star_soup = BeautifulSoup(my_page.content)
```

In [33]:

```
1 # check your code - you should see an HTML page
2 print (star_soup.prettify()[:])
```

```
<!DOCTYPE html>
<html xmlns:fb="http://www.facebook.com/2008/fbml" xmlns:og="http://
ogp.me/ns#">
  <head>
    <meta charset="utf-8"/>
    <meta content="IE=edge" http-equiv="X-UA-Compatible"/>
    <meta content="app-id=342792525, app-argument=imdb:///list/ls02581
4950?src=mdot" name="apple-itunes-app"/>
    <script type="text/javascript">
      var IMDbTimer={starttime: new Date().getTime(),pt:'java'};
    </script>
    <script>
      if (typeof uet == 'function') {
        uet("bb", "LoadTitle", {wb: 1});
      }
    </script>
    <script>
      (function(t){ (t.events = t.events || {})[ "csm_head_pre_title" ] =
new Date().getTime(); })(IMDbTimer);
    </script>
```

1.3

Function

parse_stars

Input

star_soup: the soup object with the scraped page

Returns

a list of dictionaries; each dictionary corresponds to a star profile and has the following data:

name: the name of the actor/actress as it appears at the top
gender: 0 or 1: translate the word 'actress' into 1 and 'actor' into '0'

url: the url of the link under their name that leads to a page with details

page: the string containing the soup of the text in their individual info page (from url)

Example:

```
{ 'name': 'Tom Hardy',
  'gender': 0,
  'url': 'https://www.imdb.com/name/nm0362766/?ref_=nm1ls_hd',
```

```
'page': BS object with 'html text acquired by scraping the 'url' page'  
}
```

In [176]:

```
1 # your code here  
2 import time  
3 from tqdm import tqdm  
4  
5 def parse_stars(star_soup):  
6     actors = star_soup.find_all('div', {'class': 'lister-item mode-detail'})  
7     starlist = []  
8     for actor in tqdm(actors):  
9         star = {}  
10        name = re.findall(  
11            '<a href="/name/.">(.\\w.*)',  
12            str(actor)  
13        )[0].strip()  
14        star.update({'name': name})  
15        gender = 1  
16        if 'Actor' in str(actor.find_all('p')[0]):  
17            gender = 0  
18        star.update({'gender': gender})  
19        url = 'https://www.imdb.com'  
20        url += re.findall('<a href="/name/.">\\.\\w\\.\\*', str(actor))[0]  
21        star.update({'url': url})  
22        page = requests.get(url)  
23        time.sleep(np.random.randint(8, 12))  
24        page = BeautifulSoup(page.content)  
25        star.update({'page': page})  
26        starlist.append(star)  
27    return starlist  
28  
29 starlist = parse_stars(star_soup)
```

100%|██████████| 100/100 [19:38<00:00, 11.79s/it]

In [185]:

```
1 # this list is large because of the html code into the `page` field
2 # to get a better picture, print only the first element
3 starlist[0]
```

Out[185]:

```
{'name': 'Gal Gadot',
 'gender': 1,
 'url': 'https://www.imdb.com/name/nm2933757',
 'page': <!DOCTYPE html>
<html xmlns:fb="http://www.facebook.com/2008/fbml" xmlns:og="http://ogp.me/ns#">
<head>
<meta charset="utf-8"/>
<meta content="IE=edge" http-equiv="X-UA-Compatible"/>
<meta content="app-id=342792525, app-argument=imdb:///name/nm2933757?src=mdot" name="apple-itunes-app"/>
<script type="text/javascript">var IMDbTimer={starttime: new Date().getTime(),pt:'java'};</script>
<script>
    if (typeof uet == 'function') {
        uet("bb", "LoadTitle", {wb: 1});
    }
</script>
```

Your output should look like this:

```
{'name': 'Gal Gadot',
 'gender': 1,
 'url': 'https://www.imdb.com/name/nm2933757?ref_=nmls_hd',
 'page':
<!DOCTYPE html>

<html xmlns:fb="http://www.facebook.com/2008/fbml" xmlns:og="http://ogp.me/ns#">
<head>
<meta charset="utf-8"/>
<meta content="IE=edge" http-equiv="X-UA-Compatible"/>
<meta content="app-id=342792525, app-argument=imdb:///name/nm2933757?src=mdot" name="apple-itunes-app"/>
<script type="text/javascript">var IMDbTimer={starttime: new Date().getTime(),pt:'java'};</script>
<script>
    if (typeof uet == 'function') {
        uet("bb", "LoadTitle", {wb: 1});
    }
</script>
<script>(function(t){ (t.events = t.events || {})[ "csm_head_pre_title" ] = new Date().getTime(); })(IMDbTimer);</script>

...
```

1.4

Function

create_star_table

Input

the starlist

Returns

a list of dictionaries; each dictionary corresponds to a star profile and has the following data:

star_name: the name of the actor/actress as it appears at the top

gender: 0 or 1 (1 for 'actress' and 0 for 'actor')

year_born : year they were born

first_movie: title of their first movie or TV show

year_first_movie: the year they made their first movie or TV show

credits: number of movies or TV shows they have made in their career.

Example:

```
{'star_name': 'Tom Hardy',  
  'gender': 0,  
  'year_born': 1997,  
  'first_movie' : 'Batman',  
  'year_first_movie' : 2017,  
  'credits' : 24}
```

In [217]:

```
1 # your code here
2
3 def create_star_table(starlist: list) -> list:
4     star_table = []
5     for actor in tqdm(starlist):
6         star = {}
7         star.update({'star_name': actor['name']})
8         star.update({'gender': actor['gender']})
9         search = re.findall('birth_year=(\d{4})', str(actor['page']))
10        if len(search) >= 1:
11            year_born = search[0]
12        else:
13            year_born = None
14        star.update({'year_born': year_born})
15        first_job = actor['page'].find(
16            'div',
17            {'class': 'filmo-category-section'})
18        ).find_all(
19            'div',
20            {'class': ['filmo-row even', 'filmo-row odd']})
21        )[-1]
22        first_movie = re.findall(
23            '<a href="/title/.+">(.)</a>',
24            str(first_job)
25        )[0]
26        star.update({'first_movie': first_movie})
27        first_movie_year = re.findall(
28            '.*(\d{4}).*',
29            first_job.find_all('span', {'class': 'year_column'})[0].getText()
30        )[0]
31        star.update({'year_first_movie': first_movie_year})
32        credits = re.findall('(\d+)\scredits?', str(actor['page']))[0]
33        star.update({'credits': credits})
34        star_table.append(star)
35    return star_table
```

In [218]:

```
1 # RUN THIS CELL ONLY ONCE - IT WILL TAKE SOME TIME TO RUN
2 star_table = []
3 star_table = create_star_table(starlist)
```

100%|██████████| 100/100 [00:16<00:00, 6.10it/s]

In [219]:

```
1 # check your code
2 star_table
```

Out[219]:

```
[{'star_name': 'Gal Gadot',
  'gender': 1,
  'year_born': '1985',
  'first_movie': 'Shemesh',
  'year_first_movie': '1999',
  'credits': '32'},
 {'star_name': 'Tom Hardy',
  'gender': 0,
  'year_born': '1977',
  'first_movie': 'Tommaso',
  'year_first_movie': '2001',
  'credits': '56'},
 {'star_name': 'Emilia Clarke',
  'gender': 1,
  'year_born': '1986',
  'first_movie': 'Doctors',
  'year_first_movie': '2009',
  'credits': '20'}]
```

Your output should look like this:

```
[{'name': 'Gal Gadot',
  'gender': 1,
  'year_born': '1985',
  'first_movie': 'Bubot',
  'year_first_movie': '2007',
  'credits': '25'},
 {'name': 'Tom Hardy',
  'gender': 0,
  'year_born': '1977',
  'first_movie': 'Tommaso',
  'year_first_movie': '2001',
  'credits': '55'},
 ...]
```

1.5

In [220]:

```
1 # your code here
2
3 import json
4
5 with open('star_table.json', 'w') as f:
6     json.dump(star_table, f)
```

1.6

In [243]:

```
1 # your code here
2 df = pd.read_json('data/staff_starinfo.json')
3 df.sample(10)
```

Out[243]:

	name	gender	year_born	first_movie	year_first_movie	credits
79	Joan Crawford	1	1906	Lady of the Night	1925	105
63	Daisy Ridley	1	1992	Only Yesterday	1991	32
72	Finn Jones	0	1988	Hollyoaks Later	2009	14
15	Jennifer Lawrence	1	1990	Monk	2006	30
35	Eiza González	1	1990	Lola: Érase una vez	2007	19
33	Dafne Keen	1	1966	The Refugees	2014-2015	4
43	Bill Paxton	0	1955	Crazy Mama	1975	94
61	Kaya Scodelario	1	1992	Moon	2009	23
12	Felicity Jones	1	1983	The Treasure Seekers	1998	42
58	Travis Fimmel	0	1979	I Used the Staff Solution	2001	26

In [244]:

```
1 # Check the values
2 df.describe()
```

Out[244]:

	gender	year_born	credits
count	100.000000	100.000000	100.000000
mean	0.730000	1983.080000	38.470000
std	0.446196	12.664816	22.416379
min	0.000000	1906.000000	4.000000
25%	0.000000	1979.000000	22.000000
50%	1.000000	1985.500000	34.500000
75%	1.000000	1990.000000	51.000000
max	1.000000	2004.000000	122.000000

In [245]:

```
1 # Check possible outlier
2 df[df.year_born == 1906]
```

Out[245]:

	name	gender	year_born	first_movie	year_first_movie	credits
79	Joan Crawford	1	1906	Lady of the Night	1925	105

In [246]:

```
1 # Check variables types
2 df.dtypes
```

Out[246]:

```
name          object
gender         int64
year_born      int64
first_movie    object
year_first_movie object
credits        int64
dtype: object
```

In [247]:

```
1 # Check the movies with 'year_first_movie' duplicated
2 df[df.year_first_movie.str.len() > 4]
```

Out[247]:

	name	gender	year_born	first_movie	year_first_movie	credits
3	Alexandra Daddario	1	1986	All My Children	2002-2003	51
23	Cara Delevingne	1	1992	Anna Karenina	2012/I	20
29	Robin Wright	1	1966	The Yellow Rose	1983-1984	54
33	Dafne Keen	1	1966	The Refugees	2014-2015	4
46	Jason Momoa	0	1979	Baywatch	1999-2001	27
60	Cole Sprouse	0	1992	Grace Under Fire	1993-1998	35
77	Rebecca Ferguson	1	1983	Nya tider	1999-2000	25
78	Julia Garner	1	1994	The Dreamer	2010/II	27
86	Auli'i Cravalho	1	2000	Moana	2016/I	4
96	Elodie Yung	1	1981	La vie devant nous	2002-2003	22

In [248]:

```
1 # Deal with these years
2 df.year_first_movie = pd.to_numeric(df.year_first_movie.str[:4])
3 df.sample(10)
```

Out[248]:

	name	gender	year_born	first_movie	year_first_movie	credits
93	Sophie Turner	1	1996	Another Me	2013	13
7	Dan Stevens	0	1982	Frankenstein	2004	37
59	Charlize Theron	1	1975	Children of the Corn III: Urban Harvest	1995	58
54	Katherine Waterston	1	1980	Americana	2004	43
21	Sophia Lillis	1	2002	The Lipstick Stain	2014	13
58	Travis Fimmel	0	1979	I Used the Staff Solution	2001	26
43	Bill Paxton	0	1955	Crazy Mama	1975	94
64	Emily Browning	1	1988	The Echo of Thunder	1998	30
22	Jessica Henwick	1	1992	St Trinian's 2: The Legend of Fritton's Gold	2009	25
82	Bryce Dallas Howard	1	1981	Parenthood	1989	33

In [249]:

```
1 # Check types again
2 df.dtypes
```

Out[249]:

```
name          object
gender        int64
year_born     int64
first_movie    object
year_first_movie int64
credits       int64
dtype: object
```

In [250]:

```
1 # Create new variable
2 df['age_at_first_movie'] = df.year_first_movie - df.year_born
3 df.sample(10)
```

Out[250]:

	name	gender	year_born	first_movie	year_first_movie	credits	age_at_first_movie
36	Laura Haddock	1	1985	My Family	2007	35	22
85	Connie Nielsen	1	1965	How Did You Get In? We Didn't See You Leave	1984	51	19
59	Charlize Theron	1	1975	Children of the Corn III: Urban Harvest	1995	58	20
13	Emma Stone	1	1988	The New Partridge Family	2005	42	17
41	Katheryn Winnick	1	1977	PSI Factor: Chronicles of the Paranormal	1999	61	22
99	Christian Navarro	0	1966	Day of the Dead 2: Contagium	2005	13	39
4	Bill Skarsgård	0	1990	Järngänget	2000	30	10
98	Nina Dobrev	1	1989	Playing House	2006	41	17
28	Alison Brie	1	1982	Stolen Poem	2004	61	22
68	Hugh Jackman	0	1968	Law of the Land	1994	57	26

1.7.1

In [259]:

```
1 # your code here
2 print(str(len(df[df.age_at_first_movie == 17].index))
3       + ' performers made their first movie at 17')
```

8 performers made their first movie at 17

Your output should look like this:

8 performers made their first movie at 17

1.7.2

In [263]:

```
1 # your code here
2
3 print(str(len(df[df.age_at_first_movie < 12].index))
4       + ' performers started as child actors')
```

20 performers started as child actors

1.8

In [275]:

```
1 # your code here
2 df[df.credits == df.credits.max()]
```

Out[275]:

	name	gender	year_born	first_movie	year_first_movie	credits	age_at_first_movie
42	Sean Young	1	1959	Jane Austen in Manhattan	1980	122	21

The most prolific actress or actor in IMDb's list of the Top 100 Stars for 2017 is Sean Young.

In [276]:

```
1 from IPython.core.display import HTML
2 def css_styling(): styles = open("styles/cs109.css", "r").read(); return HTML(s
3 css_styling())
```

Out[276]:

In [301]:

```
1 import seaborn as sns
2 fig, ax = plt.subplots(figsize=(10, 20))
3 sns.barplot(
4     'credits',
5     'name',
6     data=df.sort_values('credits', ascending=False),
7     ax=ax
8 )
9 ax.set_title('Actress/actor vs. credits')
10 plt.show()
```

Actress/actor vs. credits



